

▶ Datasets

[] ↳ 9 cells hidden

▼ Classification Modeling on Sentiment Prediction

```
1 # Create a copy of the bitcoin price DataFrame
2 crypto_usd.head(2)
```

	time	close	high	low	open	volumefrom	volumeto
0	2023-02-19 13:00:00	24682.03	24715.82	24682.03	24707.39	903.97	22335943.28

```
1 print(crypto_usd.columns)
```

```
Index(['time', 'close', 'high', 'low', 'open', 'volumefrom', 'volumeto',
      'Date', 'Time', 'volume', 'marketcap', 'price_delta'],
      dtype='object')
```

```
1 # Create a copy of the bitcoin tweets DataFrame
2 df_tweets = tweets.copy()
3 df_tweets.head(2)
```

	user_name	user_location	user_description	user_created	user_followers
0	lrk	Vancouver, WA	lrk started investing in the stock market in 1...	2018-08-11 03:17:00	116
1	Xiang Zhang	NaN	Professional Software Engineer ðŸ’ŽðŸ’ŽðŸ’ŽCrypto ...	2011-01-11 01:37:00	42

```
1 # Merge the tweet data with the Bitcoin price data
2 tweets_df = pd.merge(df_tweets, crypto_usd, left_on='date', right_on='time', how='inner')
```

```
1 print(tweets_df.columns)
2
```

```
Index(['user_name', 'user_location', 'user_description', 'user_created',
      'user_followers', 'user_friends', 'user_favourites', 'user_verified',
      'date', 'text', 'hashtags', 'source', 'is_retweet', 'compound', 'score',
      'sentiment_level', 'polarity', 'subjectivity', 'time', 'close', 'high',
      'low', 'open', 'volumefrom', 'volumeto', 'Date', 'Time', 'volume',
      'marketcap', 'price_delta'],
      dtype='object')
```

▼ Feature Extraction

```
1 import pandas as pd
2 from sklearn.feature_extraction.text import CountVectorizer
3 from sklearn.model_selection import train_test_split
4 from sklearn.naive_bayes import MultinomialNB
```

```

5 from sklearn.metrics import accuracy_score
6 from scipy.sparse import hstack
7
8 # Feature Extraction: Unigrams
9 unigram_vectorizer = CountVectorizer(ngram_range=(1, 1))
10 unigram_features = unigram_vectorizer.fit_transform(tweets_df['text'])
11
12 # Feature Extraction: Bigrams
13 bigram_vectorizer = CountVectorizer(ngram_range=(2, 2))
14 bigram_features = bigram_vectorizer.fit_transform(tweets_df['text'])
15
16 # Combining Features
17 combined_features = hstack([unigram_features, bigram_features])
18
19 # Perform sentiment analysis
20 X = combined_features
21 y = tweets_df['sentiment_level']
22
23 # Split the data into training and testing sets
24 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

1 #from sklearn.feature_extraction.text import CountVectorizer: This line imports the CountVectorizer class from the Scikit-learn library.

1 import numpy as np
2
3 # Print the first 10 rows of the term frequency matrix
4 print(combined_features[:10].toarray())
5

[[0 0 0 ... 0 0 0]
 [0 0 0 ... 0 0 0]
 [0 0 0 ... 0 0 0]
 ...
 [0 0 0 ... 0 0 0]
 [0 0 0 ... 0 0 0]
 [0 0 0 ... 0 0 0]]

```

▼ Naive_bayes

```

1 from sklearn.metrics import classification_report

1 from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
2
3 # Train a classification model (e.g., Naive Bayes)
4 classifier = MultinomialNB()
5 classifier.fit(X_train, y_train)
6
7 # Predict sentiment labels for test data
8 y_pred = classifier.predict(X_test)
9
10 # Evaluate the model using additional metrics
11 accuracy = accuracy_score(y_test, y_pred)
12 precision = precision_score(y_test, y_pred, average='weighted')
13 recall = recall_score(y_test, y_pred, average='weighted')
14 f1 = f1_score(y_test, y_pred, average='weighted')
15
16 print("Accuracy:", accuracy)
17 print("Precision:", precision)
18 print("Recall:", recall)
19 print("F1-Score:", f1)
20
21 # Use the trained model for future predictions
22 new_tweet = ["New tweet about Bitcoin"]
23 new_tweet_features = hstack([unigram_vectorizer.transform(new_tweet), bigram_vectorizer.transform(new_tweet)])
24 predicted_sentiment = classifier.predict(new_tweet_features)
25 #Classification Report
26 print("Predicted sentiment:", predicted_sentiment)
27 print(classification_report(y_test, y_pred))

Accuracy: 0.7879746835443038
Precision: 0.8031951087547753
Recall: 0.7879746835443038
F1-Score: 0.791715079259514

```

```

Predicted sentiment: ['Neutral']
      precision    recall  f1-score   support

Extreme Negative      0.93      0.71      0.80        55
Extreme Positive      0.50      0.69      0.58       127
      Negative      0.83      0.61      0.71       157
      Neutral      0.88      0.86      0.87       908
      Positive      0.67      0.74      0.70       333

accuracy              0.79       1580
macro avg             0.76      0.72      0.73       1580
weighted avg          0.80      0.79      0.79       1580

```

1 #The Naive Bayes classifier achieved an accuracy of 0.7879746835443038 and a precision of 0.8031951087547753 for sentiment analysis on th

▼ Support Vector Machines (SVM)

```

1 import pandas as pd
2 from sklearn.feature_extraction.text import CountVectorizer
3 from sklearn.model_selection import train_test_split
4 from sklearn.svm import LinearSVC
5 from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
6
7 try:
8     # Split the data into training and testing sets
9     X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
10
11     # Train a linear SVM classifier
12     classifier = LinearSVC()
13     classifier.fit(X_train, y_train)
14
15     # Evaluate the model using additional metrics
16     y_pred = classifier.predict(X_test)
17     accuracy = accuracy_score(y_test, y_pred)
18     precision = precision_score(y_test, y_pred, average='weighted')
19     recall = recall_score(y_test, y_pred, average='weighted')
20     f1 = f1_score(y_test, y_pred, average='weighted')
21
22     print("Accuracy:", accuracy)
23     print("Precision:", precision)
24     print("Recall:", recall)
25     print("F1-Score:", f1)
26
27     # Use the trained model for future predictions
28     new_tweet = ["New tweet about Bitcoin"]
29     new_tweet_features = hstack([unigram_vectorizer.transform(new_tweet), bigram_vectorizer.transform(new_tweet)])
30     predicted_sentiment = classifier.predict(new_tweet_features)
31     print("Predicted sentiment:", predicted_sentiment)
32
33 except Exception as e:
34     print("An error occurred:", str(e))
35 #Classification Report
36 print(classification_report(y_test, y_pred))

```

```

Accuracy: 0.8645569620253165
Precision: 0.8642113824767497
Recall: 0.8645569620253165
F1-Score: 0.859632616414193
Predicted sentiment: ['Neutral']
      precision    recall  f1-score   support

Extreme Negative      0.95      0.76      0.85        55
Extreme Positive      0.86      0.65      0.74       127
      Negative      0.89      0.70      0.78       157
      Neutral      0.87      0.97      0.92       908
      Positive      0.82      0.74      0.78       333

accuracy              0.86       1580
macro avg             0.88      0.77      0.81       1580
weighted avg          0.86      0.86      0.86       1580

```

1 #The SVM classifier achieved an accuracy of 0.8645569620253165 and a precision of 0.8642113824767497 for sentiment analysis on the tweet

▼ Random Forest

```

1 import pandas as pd
2 from sklearn.feature_extraction.text import CountVectorizer
3 from sklearn.model_selection import train_test_split
4 from sklearn.ensemble import RandomForestClassifier
5 from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
6
7
8
9 # Feature Extraction: Unigrams
10 vectorizer = CountVectorizer(ngram_range=(1, 1))
11 X = vectorizer.fit_transform(tweets_df['text'])
12 y = tweets_df['sentiment_level']
13
14 try:
15     # Split the data into training and testing sets
16     X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
17
18     # Train a linear Random Forest classifier
19     classifier = RandomForestClassifier()
20     classifier.fit(X_train, y_train)
21
22     # Evaluate the model using additional metrics
23     y_pred = classifier.predict(X_test)
24     accuracy = accuracy_score(y_test, y_pred)
25     precision = precision_score(y_test, y_pred, average='weighted')
26     recall = recall_score(y_test, y_pred, average='weighted')
27     f1 = f1_score(y_test, y_pred, average='weighted')
28
29     print("Accuracy:", accuracy)
30     print("Precision:", precision)
31     print("Recall:", recall)
32     print("F1-Score:", f1)
33
34     # Use the trained model for future predictions
35     new_tweet = ["New tweet about Bitcoin"]
36     new_tweet_features = vectorizer.transform(new_tweet)
37     predicted_sentiment = classifier.predict(new_tweet_features)
38     print("Predicted sentiment:", predicted_sentiment)
39
40 except Exception as e:
41     print("An error occurred:", str(e))
42 #Classification Report
43 print(classification_report(y_test, y_pred))

```

```

Accuracy: 0.8575949367088608
Precision: 0.8647362580034188
Recall: 0.8575949367088608
F1-Score: 0.8488711984107943
Predicted sentiment: ['Neutral']

```

	precision	recall	f1-score	support
Extreme Negative	1.00	0.75	0.85	55
Extreme Positive	0.94	0.47	0.63	127
Negative	0.95	0.66	0.78	157
Neutral	0.85	0.99	0.91	908
Positive	0.81	0.77	0.79	333
accuracy			0.86	1580
macro avg	0.91	0.73	0.79	1580
weighted avg	0.86	0.86	0.85	1580

1 #The Random Forest classifier achieved an accuracy of 0.870253164556962 and a precision of 0.8670628029958947 for sentiment analysis on t

▼ Logistic Regression

```

1 import pandas as pd
2 from sklearn.feature_extraction.text import CountVectorizer
3 from sklearn.model_selection import train_test_split
4 from sklearn.linear_model import LogisticRegression

```

```

5 from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
6
7 # Assuming you have tweets_df with the appropriate 'text' and 'sentiment_level' columns
8
9 # Feature Extraction: Unigrams
10 vectorizer = CountVectorizer(ngram_range=(1, 1))
11 X = vectorizer.fit_transform(tweets_df['text'])
12 y = tweets_df['sentiment_level']
13
14 try:
15     # Split the data into training and testing sets
16     X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
17
18     # Train a logistic regression classifier with increased max_iter
19     classifier = LogisticRegression(max_iter=1000)
20     classifier.fit(X_train, y_train)
21
22     # Evaluate the model using additional metrics
23     y_pred = classifier.predict(X_test)
24     accuracy = accuracy_score(y_test, y_pred)
25     precision = precision_score(y_test, y_pred, average='weighted')
26     recall = recall_score(y_test, y_pred, average='weighted')
27     f1 = f1_score(y_test, y_pred, average='weighted')
28
29     print("Accuracy:", accuracy)
30     print("Precision:", precision)
31     print("Recall:", recall)
32     print("F1-Score:", f1)
33
34     # Use the trained model for future predictions
35     new_tweet = ["New tweet about Bitcoin"]
36     new_tweet_features = vectorizer.transform(new_tweet)
37     predicted_sentiment = classifier.predict(new_tweet_features)
38     print("Predicted sentiment:", predicted_sentiment)
39
40 except Exception as e:
41     print("An error occurred:", str(e))
42 #Classification Report
43 print(classification_report(y_test, y_pred))

```

```

Accuracy: 0.859493670886076
Precision: 0.8596313804881421
Recall: 0.859493670886076
F1-Score: 0.8545443425051455
Predicted sentiment: ['Neutral']

```

	precision	recall	f1-score	support
Extreme Negative	1.00	0.73	0.84	55
Extreme Positive	0.86	0.64	0.73	127
Negative	0.86	0.69	0.76	157
Neutral	0.87	0.97	0.91	908
Positive	0.82	0.75	0.78	333
accuracy			0.86	1580
macro avg	0.88	0.75	0.81	1580
weighted avg	0.86	0.86	0.85	1580

1 #The Logistic Regression classifier achieved an accuracy of 0.859493670886076 and a precision of 0.8596313804881421 for sentiment analysis

▼ Gradient Boosting

```

1 import pandas as pd
2 from sklearn.feature_extraction.text import CountVectorizer
3 from sklearn.model_selection import train_test_split
4 from sklearn.ensemble import GradientBoostingClassifier
5 from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
6
7 # Assuming you have tweets_df with the appropriate 'text' and 'sentiment_level' columns
8
9 # Feature Extraction: Unigrams
10 vectorizer = CountVectorizer(ngram_range=(1, 1))
11 X = vectorizer.fit_transform(tweets_df['text'])
12 y = tweets_df['sentiment_level']
13

```

```

14 try:
15     # Split the data into training and testing sets
16     X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
17
18     # Train a Gradient Boosting classifier
19     classifier = GradientBoostingClassifier()
20     classifier.fit(X_train, y_train)
21
22     # Evaluate the model using additional metrics
23     y_pred = classifier.predict(X_test)
24     accuracy = accuracy_score(y_test, y_pred)
25     precision = precision_score(y_test, y_pred, average='weighted')
26     recall = recall_score(y_test, y_pred, average='weighted')
27     f1 = f1_score(y_test, y_pred, average='weighted')
28
29     print("Accuracy:", accuracy)
30     print("Precision:", precision)
31     print("Recall:", recall)
32     print("F1-Score:", f1)
33
34     # Use the trained model for future predictions
35     new_tweet = ["New tweet about Bitcoin"]
36     new_tweet_features = vectorizer.transform(new_tweet)
37     predicted_sentiment = classifier.predict(new_tweet_features)
38     print("Predicted sentiment:", predicted_sentiment)
39
40 except Exception as e:
41     print("An error occurred:", str(e))
42 #Classification Report
43 print(classification_report(y_test, y_pred))

```

```

Accuracy: 0.8468354430379746
Precision: 0.8543684218834472
Recall: 0.8468354430379746
F1-Score: 0.8375792142254445
Predicted sentiment: ['Neutral']

```

	precision	recall	f1-score	support
Extreme Negative	0.90	0.78	0.83	55
Extreme Positive	0.94	0.57	0.71	127
Negative	0.92	0.62	0.74	157
Neutral	0.83	0.99	0.90	908
Positive	0.86	0.68	0.76	333
accuracy			0.85	1580
macro avg	0.89	0.73	0.79	1580
weighted avg	0.85	0.85	0.84	1580

1 #The Gradient Boosting classifier achieved an accuracy of 0.8468354430379746 and a precision of 0.8543684218834472 for sentiment analysis

▼ Cross Validation of Models

```

1 from sklearn.naive_bayes import MultinomialNB
2 from sklearn.svm import SVC
3 from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
4 from sklearn.linear_model import LogisticRegression
5 from sklearn.model_selection import cross_val_score
6
7 # Define the models
8 models = [
9     ("Naive Bayes", MultinomialNB()),
10    ("Support Vector Machine", SVC()),
11    ("Random Forest", RandomForestClassifier()),
12    ("Logistic Regression", LogisticRegression()),
13    ("Gradient Boosting", GradientBoostingClassifier())
14 ]
15
16 # Perform cross-validation and evaluation for each model
17 for model_name, model in models:
18     # Perform cross-validation
19     scores = cross_val_score(model, X_train, y_train, cv=5)
20     mean_score = scores.mean()
21

```

```

22 # Fit the model on the entire training set
23 model.fit(X_train, y_train)
24
25 # Evaluate the model on the test set
26 accuracy = model.score(X_test, y_test)
27
28 # Print the results
29 print("Model:", model_name)
30 print("Cross-Validation Mean Score:", mean_score)
31 print("Accuracy:", accuracy)
32 print()
33

```

```

Model: Naive Bayes
Cross-Validation Mean Score: 0.7910727171592651
Accuracy: 0.7879746835443038

```

▼ Cross Validation

```
1 #cross-validation for the models using scikit-learn's cross_val_score function
```

```

1 import pandas as pd
2 from sklearn.feature_extraction.text import CountVectorizer
3 from sklearn.feature_selection import SelectKBest, chi2
4 from sklearn.model_selection import train_test_split, cross_val_score
5 from sklearn.naive_bayes import MultinomialNB
6 from sklearn.svm import SVC
7 from sklearn.ensemble import RandomForestClassifier
8 from sklearn.metrics import accuracy_score
9 from scipy.sparse import hstack
10

```

```
1 #Naive Bayes
```

```

1 from sklearn.naive_bayes import MultinomialNB
2 from sklearn.svm import LinearSVC
3 from sklearn.ensemble import RandomForestClassifier
4 from sklearn.model_selection import cross_val_score
5
6 # Train and evaluate Naive Bayes
7 naive_bayes = MultinomialNB()
8 naive_bayes_scores = cross_val_score(naive_bayes, X_train, y_train, cv=5)
9 print("Naive Bayes Cross-Validation Scores:", naive_bayes_scores.mean())
10 naive_bayes.fit(X_train, y_train)
11 naive_bayes_accuracy = naive_bayes.score(X_test, y_test)
12 print("Naive Bayes Accuracy:", naive_bayes_accuracy)
13 # Predict sentiment labels for test data
14 y_pred = naive_bayes.predict(X_test)
15 from sklearn.metrics import classification_report
16 print(classification_report(y_test, y_pred))

```

```

Naive Bayes Cross-Validation Scores: 0.7888584042414585
Naive Bayes Accuracy: 0.7943037974683544

```

	precision	recall	f1-score	support
Extreme Negative	0.94	0.58	0.72	55
Extreme Positive	0.60	0.57	0.59	127
Negative	0.85	0.58	0.69	157
Neutral	0.84	0.92	0.88	908
Positive	0.69	0.68	0.69	333
accuracy			0.79	1580
macro avg	0.79	0.67	0.71	1580
weighted avg	0.79	0.79	0.79	1580

```
1 #SVM
```

```

1 from sklearn.naive_bayes import MultinomialNB
2 from sklearn.svm import LinearSVC
3 from sklearn.ensemble import RandomForestClassifier

```

```

4 from sklearn.model_selection import cross_val_score
5
6 # Train and evaluate SVM
7 svm = LinearSVC()
8 svm_scores = cross_val_score(svm, X_train, y_train, cv=5)
9 print("SVM Cross-Validation Scores:", svm_scores.mean())
10 svm.fit(X_train, y_train)
11 svm_accuracy = svm.score(X_test, y_test)
12 print("SVM Accuracy:", svm_accuracy)
13 # Predict sentiment labels for test data
14 y_pred = svm.predict(X_test)
15 from sklearn.metrics import classification_report
16 print(classification_report(y_test, y_pred))

```

```

SVM Cross-Validation Scores: 0.8630914439199415
SVM Accuracy: 0.870253164556962

```

	precision	recall	f1-score	support
Extreme Negative	0.89	0.76	0.82	55
Extreme Positive	0.79	0.70	0.74	127
Negative	0.81	0.71	0.76	157
Neutral	0.90	0.96	0.93	908
Positive	0.83	0.77	0.80	333
accuracy			0.87	1580
macro avg	0.84	0.78	0.81	1580
weighted avg	0.87	0.87	0.87	1580

```
1 #Random Forest
```

```

1
2 # Train Random Forest classifier
3 random_forest = RandomForestClassifier(n_estimators=100, n_jobs=-1)
4 random_forest.fit(X_train, y_train)
5
6 # Evaluate Random Forest
7 random_forest_scores = cross_val_score(random_forest, X_train, y_train, cv=5)
8 random_forest_mean_score = random_forest_scores.mean()
9
10 random_forest_accuracy = random_forest.score(X_test, y_test)
11
12 # Print results
13 print("Random Forest Cross-Validation Mean Score:", random_forest_mean_score)
14 print("Random Forest Accuracy:", random_forest_accuracy)
15 # Predict sentiment labels for test data
16 y_pred = random_forest.predict(X_test)
17 from sklearn.metrics import classification_report
18 print(classification_report(y_test, y_pred))
19

```

```

Random Forest Cross-Validation Mean Score: 0.8553332681880594
Random Forest Accuracy: 0.8645569620253165

```

	precision	recall	f1-score	support
Extreme Negative	1.00	0.75	0.85	55
Extreme Positive	0.96	0.52	0.67	127
Negative	0.94	0.66	0.78	157
Neutral	0.85	0.99	0.92	908
Positive	0.83	0.78	0.80	333
accuracy			0.86	1580
macro avg	0.92	0.74	0.80	1580
weighted avg	0.87	0.86	0.86	1580

```
1 #Logistic Regression
```

```

1 from sklearn.naive_bayes import MultinomialNB
2 from sklearn.svm import LinearSVC
3 from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
4 from sklearn.linear_model import LogisticRegression
5 from sklearn.model_selection import cross_val_score
6
7 # Train and evaluate Logistic Regression
8 logistic_regression = LogisticRegression(max_iter=1000)

```



```

9 logistic_regression_scores = cross_val_score(logistic_regression, X_train, y_train, cv=5)
10 logistic_regression_mean_score = logistic_regression_scores.mean()
11 logistic_regression.fit(X_train, y_train)
12 logistic_regression_accuracy = logistic_regression.score(X_test, y_test)
13 print("Logistic Regression Cross-Validation Mean Score:", logistic_regression_mean_score)
14 print("Logistic Regression Accuracy:", logistic_regression_accuracy)
15 # Predict sentiment labels for test data
16 y_pred = logistic_regression.predict(X_test)
17 from sklearn.metrics import classification_report
18 print(classification_report(y_test, y_pred))

```

Logistic Regression Cross-Validation Mean Score: 0.8429882387724625
 Logistic Regression Accuracy: 0.8537974683544304

	precision	recall	f1-score	support
Extreme Negative	0.98	0.75	0.85	55
Extreme Positive	0.90	0.55	0.68	127
Negative	0.92	0.69	0.79	157
Neutral	0.84	0.98	0.91	908
Positive	0.83	0.73	0.78	333
accuracy			0.85	1580
macro avg	0.89	0.74	0.80	1580
weighted avg	0.86	0.85	0.85	1580

1 #Gradient Boosting

```

1 from sklearn.ensemble import GradientBoostingClassifier
2 from sklearn.model_selection import cross_val_score
3
4 # Train and evaluate Gradient Boosting Classifier
5 gradient_boosting = GradientBoostingClassifier()
6 gradient_boosting_scores = cross_val_score(gradient_boosting, X_train, y_train, cv=3) # Adjust cv parameter as needed
7 gradient_boosting_mean_score = gradient_boosting_scores.mean()
8
9 gradient_boosting.fit(X_train, y_train)
10 gradient_boosting_accuracy = gradient_boosting.score(X_test, y_test)
11
12 print("Gradient Boosting Cross-Validation Mean Score:", gradient_boosting_mean_score)
13 print("Gradient Boosting Accuracy:", gradient_boosting_accuracy)
14 # Predict sentiment labels for test data
15 y_pred = gradient_boosting.predict(X_test)
16 from sklearn.metrics import classification_report
17 print(classification_report(y_test, y_pred))

```

Gradient Boosting Cross-Validation Mean Score: 0.8421968977524533
 Gradient Boosting Accuracy: 0.8436708860759494

	precision	recall	f1-score	support
Extreme Negative	0.89	0.76	0.82	55
Extreme Positive	0.93	0.58	0.71	127
Negative	0.93	0.61	0.74	157
Neutral	0.82	0.99	0.90	908
Positive	0.86	0.66	0.75	333
accuracy			0.84	1580
macro avg	0.89	0.72	0.78	1580
weighted avg	0.85	0.84	0.83	1580

▼ Hyperparameter Tuning

```

1 import pandas as pd
2 from sklearn.feature_extraction.text import CountVectorizer
3 from sklearn.feature_selection import SelectKBest, chi2
4 from sklearn.model_selection import train_test_split, cross_val_score, GridSearchCV
5 from sklearn.naive_bayes import MultinomialNB
6 from sklearn.svm import SVC
7 from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
8 from sklearn.linear_model import LogisticRegression
9 from sklearn.metrics import accuracy_score
10 from scipy.sparse import hstack
11
12 # Feature Extraction: Unigrams

```

```

13 unigram_vectorizer = CountVectorizer(ngram_range=(1, 1))
14 unigram_features = unigram_vectorizer.fit_transform(tweets_df['text'])
15
16 # Feature Extraction: Bigrams
17 bigram_vectorizer = CountVectorizer(ngram_range=(2, 2))
18 bigram_features = bigram_vectorizer.fit_transform(tweets_df['text'])
19
20 # Combining Features
21 combined_features = hstack([unigram_features, bigram_features])
22
23 # Perform sentiment analysis
24 X = combined_features
25 y = tweets_df['sentiment_level']
26
27 # Apply feature selection
28 k = 1000 # Number of top features to select
29 feature_selector = SelectKBest(chi2, k=k)
30 X_selected = feature_selector.fit_transform(X, y)
31
32 # Split the data into training and testing sets
33 X_train, X_test, y_train, y_test = train_test_split(X_selected, y, test_size=0.2, random_state=42)
34
35 # Define the models and their respective hyperparameter grids
36 models = [
37     ("Naive Bayes", MultinomialNB(), {'alpha': [0.1, 1.0, 10.0]}),
38     ("Support Vector Machine", SVC(), {'C': [0.1, 1.0, 10.0]}),
39     ("Random Forest", RandomForestClassifier(), {'n_estimators': [100, 200, 300]}),
40     ("Logistic Regression", LogisticRegression(), {'C': [0.1, 1.0, 10.0]}),
41     ("Gradient Boosting", GradientBoostingClassifier(), {'n_estimators': [100, 200, 300]})
42 ]
43
44 # Perform cross-validation and evaluation for each model
45 for model_name, model, param_grid in models:
46     # Perform hyperparameter tuning using GridSearchCV
47     grid_search = GridSearchCV(model, param_grid, cv=5)
48     grid_search.fit(X_train, y_train)
49
50     # Get the best model and its parameters
51     best_model = grid_search.best_estimator_
52     best_params = grid_search.best_params_
53
54     # Perform cross-validation with the best model
55     cross_val_scores = cross_val_score(best_model, X_train, y_train, cv=5)
56
57     # Fit the best model on the entire training set
58     best_model.fit(X_train, y_train)
59
60     # Make predictions on the test set
61     y_pred = best_model.predict(X_test)
62
63     # Calculate accuracy
64     accuracy = accuracy_score(y_test, y_pred)
65
66     # Print the results
67     print("Model:", model_name)
68     print("Best Parameters:", best_params)
69     print("Cross-Validation Accuracy:", cross_val_scores.mean())
70     print("Accuracy:", accuracy)
71     print()
72

```

```

Model: Naive Bayes
Best Parameters: {'alpha': 1.0}
Cross-Validation Accuracy: 0.7736638954869359
Accuracy: 0.7639240506329114

```

```

Model: Support Vector Machine
Best Parameters: {'C': 10.0}
Cross-Validation Accuracy: 0.8809735710634715
Accuracy: 0.8848101265822785

```

```

Model: Random Forest
Best Parameters: {'n_estimators': 300}
Cross-Validation Accuracy: 0.8649886747446806
Accuracy: 0.8613924050632912

```

```

/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

▼ Model TPOT

```
1 # Assuming you have the 'data1' and 'data2' DataFrames
2 data1 = crypto_usd.copy()
3 data2 = tweets.copy()
4 # Merge the two DataFrames based on 'time' and 'date' columns
5 merge = pd.merge(data1, data2, left_on='time', right_on='date')
6
7 # Drop the duplicate 'date' column
8 merge.drop('date', axis=1, inplace=True)
9
10 # Display the merged DataFrame
11 print(merge)
12
```

		time	close	high	low	open	volume	from \
0	2023-02-25	21:00:00	22944.16	22960.69	22863.96	22921.71	1331.05	
1	2023-02-25	21:00:00	22944.16	22960.69	22863.96	22921.71	1331.05	
2	2023-02-25	21:00:00	22944.16	22960.69	22863.96	22921.71	1331.05	
3	2023-02-25	21:00:00	22944.16	22960.69	22863.96	22921.71	1331.05	
4	2023-02-25	21:00:00	22944.16	22960.69	22863.96	22921.71	1331.05	
...	
7893	2023-03-04	23:00:00	22351.08	22352.28	22302.56	22311.46	476.12	
7894	2023-03-04	23:00:00	22351.08	22352.28	22302.56	22311.46	476.12	
7895	2023-03-04	23:00:00	22351.08	22352.28	22302.56	22311.46	476.12	
7896	2023-03-04	23:00:00	22351.08	22352.28	22302.56	22311.46	476.12	
7897	2023-03-04	23:00:00	22351.08	22352.28	22302.56	22311.46	476.12	

	volume	to	Date	Time	volume	...	user_verified	\
0	30505954.61	2023-02-25	21:00:00	30504623.56	...	False		
1	30505954.61	2023-02-25	21:00:00	30504623.56	...	False		
2	30505954.61	2023-02-25	21:00:00	30504623.56	...	False		
3	30505954.61	2023-02-25	21:00:00	30504623.56	...	False		
4	30505954.61	2023-02-25	21:00:00	30504623.56	...	False		
...		
7893	10632637.83	2023-03-04	23:00:00	10632161.71	...	False		
7894	10632637.83	2023-03-04	23:00:00	10632161.71	...	False		
7895	10632637.83	2023-03-04	23:00:00	10632161.71	...	False		

```
7896 10632637.83 2023-03-04 23:00:00 10632161.71 ... False
7897 10632637.83 2023-03-04 23:00:00 10632161.71 ... False
```

```

                                text \
0  ethereum price updat eth 157128 usd bitcoin 00...
1                                bitcoin 1month predict tuhgbqkln
2  btcusdt 15m volum spike btc btc bitcoin ucl5iaaq4
3  l0000k take time think littlebit person load a...
4  000000, sat 25 feb 2023 210035 gmt top 10 btc...
...
7893  usd racist built colonist slaver paid btc bc e...
7894  everris rise everrisev3 everrevok defi crypto ...
7895  0000 parti time 0000 0000 10000 x1 megapr 0000...
7896  strategi 5010hl1h atr20d 92138 04 mar 2023 230...
7897  complet variou task hh8vl67nz5 claim slm token...
```

```

                                hashtags                                source \
0  ['Ethereum', 'ETH', 'Bitcoin', 'BTC', 'altcoin...  Twitter Web App
1                                ['Bitcoin']  predictCCbot
2                                ['BTC', 'Bitcoin']  JumplineAlerts
3  ['GGA', 'cryptocurrency', 'Bitcoin', 'bnb', 'T...  Twitter for Android
4                                ['bitcoin']  eht10c
...
7893                                ['BTC']  Twitter for Android
7894  ['EverRise', 'EverRiseV3', 'EverRevoke', 'DeFi...  EverRiseTwitterBot1
7895  ['btc', 'eth', 'xrp', 'doge', 'shiba', 'lto', ...  Twitter Web App
7896                                ['BTC', 'BitMEX']  system'cRe5520'
7897  ['SLMGames', 'SLM', 'Web3', 'BTC', 'ETH', 'BSC...  TweetDeck
```


```

is_retweet compound score sentiment_level polarity \
0  0.0 0.0000 0.000000e+00 Neutral 0.000000
1  0.0 0.0000 0.000000e+00 Neutral 0.000000
2  0.0 0.0000 0.000000e+00 Neutral 0.000000
3  0.0 -0.3089 -9.666133e+05 Negative -0.041667
4  0.0 0.2023 7.485100e+00 Positive 0.500000
```

1 merge.head()

	time	close	high	low	open	volumefrom	volumeto	Date	Time	volume	...	user_verified	text
0	2023-02-25 21:00:00	22944.16	22960.69	22863.96	22921.71	1331.05	30505954.61	2023-02-25	21:00:00	30504623.56	...	False	ethereum price updat eth 157128 usd bitcoin 00...
1	2023-02-25 21:00:00	22944.16	22960.69	22863.96	22921.71	1331.05	30505954.61	2023-02-25	21:00:00	30504623.56	...	False	bitcoin 1month predict tuhgbqkln
2	2023-02-25 21:00:00	22944.16	22960.69	22863.96	22921.71	1331.05	30505954.61	2023-02-25	21:00:00	30504623.56	...	False	btcusdt 15m volum spike btc btc bitcoin ucl5iaaq4
3	2023-02-25 21:00:00	22944.16	22960.69	22863.96	22921.71	1331.05	30505954.61	2023-02-25	21:00:00	30504623.56	...	False	l0000k take time think littlebit person load a...
4	2023-02-25 21:00:00	22944.16	22960.69	22863.96	22921.71	1331.05	30505954.61	2023-02-25	21:00:00	30504623.56	...	False	000000, sat 25 feb 2023 210035 gmt top 10 btc...

5 rows × 29 columns



1 merge.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 7898 entries, 0 to 7897
Data columns (total 29 columns):
#   Column                Non-Null Count  Dtype
...
```

```

---  -----
0   time           7898 non-null object
1   close          7898 non-null float64
2   high           7898 non-null float64
3   low            7898 non-null float64
4   open           7898 non-null float64
5   volumefrom     7898 non-null float64
6   volumeto       7898 non-null float64
7   Date           7898 non-null object
8   Time           7898 non-null object
9   volume         7898 non-null float64
10  marketcap      7898 non-null float64
11  price_delta    7898 non-null float64
12  user_name      7898 non-null object
13  user_location  3898 non-null object
14  user_description 7620 non-null object
15  user_created   7898 non-null object
16  user_followers 7898 non-null float64
17  user_friends   7898 non-null float64
18  user_favourites 7898 non-null float64
19  user_verified  7898 non-null bool
20  text           7898 non-null object
21  hashtags       7891 non-null object
22  source         7891 non-null object
23  is_retweet     7891 non-null float64
24  compound        7898 non-null float64
25  score          7898 non-null float64
26  sentiment_level 7898 non-null object
27  polarity        7898 non-null float64
28  subjectivity   7898 non-null float64
dtypes: bool(1), float64(17), object(11)
memory usage: 1.8+ MB

```

```

1 label_counts = tweets['sentiment_level'].value_counts()
2 print(label_counts)

```

```

Neutral          93169
Positive         35921
Extreme Positive  17343
Negative         15903
Extreme Negative  5316
Name: sentiment_level, dtype: int64


```

```

1 import matplotlib.pyplot as plt
2 # scatter plot to show the subjectivity and the polarity
3 plt.figure(figsize=(14,10))
4
5 for i in range(merge.shape[0]):
6     plt.scatter(merge["polarity"].iloc[[i]].values[0], merge["subjectivity"].iloc[[i]].values[0], color="Purple")
7
8 plt.title("Sentiment Analysis Scatter Plot")
9 plt.xlabel('polarity')
10 plt.ylabel('subjectivity')
11 plt.show()

```

Sentiment Analysis Scatter Plot



```
1 #Creating Target Column
```

```
1 price_indicator = [merge.close[0] - merge['open'][0]]
2 for i in range(99):
3     price_indicator.append(merge.close[i+1] - merge.close[i])
4 #price_indicator
```

```
1 merge['price_indicator'] = 0
2 for i in range(len(price_indicator)):
3     merge['price_indicator'][i] = price_indicator[i]
4
5 merge.head()
```

<ipython-input-18-8f90b0759c32>:3: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

	time	close	high	low	open	volumefrom	volumeto	Date	Time	volume	...	text	hashtags
0	2023-02-25 21:00:00	22944.16	22960.69	22863.96	22921.71	1331.05	30505954.61	2023-02-25	21:00:00	30504623.56	...	ethereum price updat eth 157128 usd bitcoin 00...	['Ethereum', 'ETH', 'Bitcoin', 'BTC', 'altcoin...']
1	2023-02-25 21:00:00	22944.16	22960.69	22863.96	22921.71	1331.05	30505954.61	2023-02-25	21:00:00	30504623.56	...	bitcoin 1month predict tuhgbqklxn	['Bitcoin']
2	2023-02-25 21:00:00	22944.16	22960.69	22863.96	22921.71	1331.05	30505954.61	2023-02-25	21:00:00	30504623.56	...	btcsdt 15m volum spike btc btc bitcoin ucl5iaaq4	['BTC', 'Bitcoin']
3	2023-02-25 21:00:00	22944.16	22960.69	22863.96	22921.71	1331.05	30505954.61	2023-02-25	21:00:00	30504623.56	...	l0k take time think littlebit person load a...	['GGA', 'cryptocurrency', 'Bitcoin', 'bnb', 'T...']
4	2023-02-25 21:00:00	22944.16	22960.69	22863.96	22921.71	1331.05	30505954.61	2023-02-25	21:00:00	30504623.56	...	δ sat 25 feb 2023 210035 gmt top 10 btc...	['bitcoin']

5 rows × 30 columns

```

1 merge['target'] = 0
2 for i in range(100):
3     if merge.price_indicator[i] > 0:
4         merge['target'][i] = 1
5
6 # 0 - price down
7 # 1 - price up
8
9 merge.head()

```

<ipython-input-19-e0c87f2219f9>:4: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-merge['target'][i] = 1)

	time	close	high	low	open	volumefrom	volumeto	Date	Time	volume	...	hashtags	source
0	2023-02-25 21:00:00	22944.16	22960.69	22863.96	22921.71	1331.05	30505954.61	2023-02-25	21:00:00	30504623.56	...	['Ethereum', 'ETH', 'Bitcoin', 'BTC', 'altcoin...	Twitter Wel Apj
1	2023-02-25 21:00:00	22944.16	22960.69	22863.96	22921.71	1331.05	30505954.61	2023-02-25	21:00:00	30504623.56	...	['Bitcoin']	predictCCbo
2	2023-02-25 21:00:00	22944.16	22960.69	22863.96	22921.71	1331.05	30505954.61	2023-02-25	21:00:00	30504623.56	...	['BTC', 'Bitcoin']	JumpLineAlert
3	2023-02-25 21:00:00	22944.16	22960.69	22863.96	22921.71	1331.05	30505954.61	2023-02-25	21:00:00	30504623.56	...	['GGA', 'cryptocurrency', 'Bitcoin', 'bnb', 'T...	Twitter fo Androi
4	2023-02-25 21:00:00	22944.16	22960.69	22863.96	22921.71	1331.05	30505954.61	2023-02-25	21:00:00	30504623.56	...	['bitcoin']	eht10r

5 rows × 31 columns

```

1 keep_columns = ['open', 'high', 'low', 'close', 'volume', 'polarity', 'subjectivity', 'compound', 'score', 'price_indicator', 'target']
2 df = merge[keep_columns]
3 df.head()

```

	open	high	low	close	volume	polarity	subjectivity	compound	score	price_indicator	target
0	22921.71	22960.69	22863.96	22944.16	30504623.56	0.000000	0.250000	0.0000	0.0000	22.45	1
1	22921.71	22960.69	22863.96	22944.16	30504623.56	0.000000	0.000000	0.0000	0.0000	0.00	0
2	22921.71	22960.69	22863.96	22944.16	30504623.56	0.000000	0.000000	0.0000	0.0000	0.00	0
3	22921.71	22960.69	22863.96	22944.16	30504623.56	-0.041667	0.458333	-0.3089	-966613.2779	0.00	0
4	22921.71	22960.69	22863.96	22944.16	30504623.56	0.500000	0.500000	0.2023	7.4851	0.00	0

1 #Model Building

```

1 import numpy as np
2 #Create the feature data set
3 X = df
4 X = np.array(X.drop(['target'],1))
5 #Create the target data set
6 y = np.array(df['target'])

```

<ipython-input-22-63d9de6a3c5f>:4: FutureWarning: In a future version of pandas all arguments of DataFrame.drop except for the argument X = np.array(X.drop(['target'],1))

```

1 from sklearn.model_selection import train_test_split
2 #Split the data into 80% training and 20% testing data sets
3 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state = 0)

```

```

1 !pip install tpot
2

Collecting tpot
  Downloading TPOT-0.12.0-py3-none-any.whl (87 kB)
    87.4/87.4 kB 5.8 MB/s eta 0:00:00
Requirement already satisfied: numpy>=1.16.3 in /usr/local/lib/python3.10/dist-packages (from tpot) (1.22.4)
Requirement already satisfied: scipy>=1.3.1 in /usr/local/lib/python3.10/dist-packages (from tpot) (1.10.1)
Requirement already satisfied: scikit-learn>=0.22.0 in /usr/local/lib/python3.10/dist-packages (from tpot) (1.2.2)
Collecting deap>=1.2 (from tpot)
  Downloading deap-1.3.3-cp310-cp310-manylinux_2_5_x86_64.manylinux1_x86_64.manylinux_2_17_x86_64.manylinux2014_x86_64.whl (139 kB)
    139.9/139.9 kB 14.0 MB/s eta 0:00:00
Collecting update-checker>=0.16 (from tpot)
  Downloading update_checker-0.18.0-py3-none-any.whl (7.0 kB)
Requirement already satisfied: tqdm>=4.36.1 in /usr/local/lib/python3.10/dist-packages (from tpot) (4.65.0)
Collecting stopit>=1.1.1 (from tpot)
  Downloading stopit-1.1.2.tar.gz (18 kB)
  Preparing metadata (setup.py) ... done
Requirement already satisfied: pandas>=0.24.2 in /usr/local/lib/python3.10/dist-packages (from tpot) (1.5.3)
Requirement already satisfied: joblib>=0.13.2 in /usr/local/lib/python3.10/dist-packages (from tpot) (1.2.0)
Requirement already satisfied: xgboost>=1.1.0 in /usr/local/lib/python3.10/dist-packages (from tpot) (1.7.6)
Requirement already satisfied: python-dateutil>=2.8.1 in /usr/local/lib/python3.10/dist-packages (from pandas>=0.24.2->tpot) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas>=0.24.2->tpot) (2022.7.1)
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn>=0.22.0->tpot) (3.1.0)
Requirement already satisfied: requests>=2.3.0 in /usr/local/lib/python3.10/dist-packages (from update-checker>=0.16->tpot) (2.27.1)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.8.1->pandas>=0.24.2->tpot) (1.16.0)
Requirement already satisfied: urllib3<1.27,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests>=2.3.0->update-checker>=0.16->tpot) (1.26.15)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests>=2.3.0->update-checker>=0.16->tpot) (2022.9.24)
Requirement already satisfied: charset-normalizer~=2.0.0 in /usr/local/lib/python3.10/dist-packages (from requests>=2.3.0->update-checker>=0.16->tpot) (2.0.12)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests>=2.3.0->update-checker>=0.16->tpot) (3.4)
Building wheels for collected packages: stopit
  Building wheel for stopit (setup.py) ... done
  Created wheel for stopit: filename=stopit-1.1.2-py3-none-any.whl size=11938 sha256=a56fda5b968cc0cd8d28799e3e03a41bd1d28bf77cd34f246
  Stored in directory: /root/.cache/pip/wheels/af/f9/87/bf5b3d565c2a007b4dae9d8142dccc85a9f164e517062dd519
Successfully built stopit
Installing collected packages: stopit, deap, update-checker, tpot
Successfully installed deap-1.3.3 stopit-1.1.2 tpot-0.12.0 update-checker-0.18.0

```

```

1 from tpot import TPOTClassifier
2 from sklearn.metrics import confusion_matrix,accuracy_score,roc_auc_score

```

```

1 from sklearn.metrics import roc_auc_score
2 from tpot import TPOTClassifier
3 import numpy as np
4

```

```

1 from sklearn.metrics import roc_auc_score
2 from tpot import TPOTClassifier
3 import numpy as np
4
5 # Instantiate TPOTClassifier
6 tpot = TPOTClassifier(
7     generations=5,
8     population_size=20,
9     verbosity=2,
10    scoring='roc_auc',
11    random_state=42,
12    disable_update_check=True,
13    config_dict='TPOT light'
14 )
15
16 # Convert X_train and y_train to NumPy arrays
17 X_train = np.array(X_train)
18 y_train = np.array(y_train)
19
20 # Ensure that there are at least two classes in y_train
21 if len(np.unique(y_train)) < 2:
22     raise ValueError("At least two classes are required in y_train for ROC AUC score calculation.")
23
24 try:
25     # Fit TPOTClassifier
26     tpot.fit(X_train, y_train)
27
28     # AUC score for tpot model
29     X_test = np.array(X_test) # Assuming you have X_test data
30     y_test = np.array(y_test) # Assuming you have y_test data

```



```

31
32 # Ensure that there are at least two classes in y_test
33 if len(np.unique(y_test)) < 2:
34     raise ValueError("At least two classes are required in y_test for ROC AUC score calculation.")
35
36 tpot_auc_score = roc_auc_score(y_test, tpot.predict_proba(X_test)[:, 1])
37 print(f'\nAUC score: {tpot_auc_score:.4f}')
38
39 # Print best pipeline steps
40 print('\nBest pipeline steps:')
41 for idx, (name, transform) in enumerate(tpot.fitted_pipeline_.steps, start=1):
42     print(f'{idx}. {transform}')
43
44 except ValueError as e:
45     print("Error:", str(e))
46

```

Optimization Progress: 33%

40/120 [00:11<00:29, 2.76pipeline/s]

```

/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_split.py:700: UserWarning: The least populated class in y has only 1
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_split.py:700: UserWarning: The least populated class in y has only 1
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_split.py:700: UserWarning: The least populated class in y has only 1
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_split.py:700: UserWarning: The least populated class in y has only 1
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_split.py:700: UserWarning: The least populated class in y has only 1
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_split.py:700: UserWarning: The least populated class in y has only 1
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_split.py:700: UserWarning: The least populated class in y has only 1
warnings.warn(
Error: Only one class present in y_true. ROC AUC score is not defined in that case.
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_split.py:700: UserWarning: The least populated class in y has only 1
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_split.py:700: UserWarning: The least populated class in y has only 1
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_split.py:700: UserWarning: The least populated class in y has only 1
warnings.warn(

```

```

1 # Instantiate TPOTClassifier
2 tpot = TPOTClassifier(
3     generations=5, #number of iterations to run ; pipeline optimisation process ; by default value is 100
4     population_size=20, #number of individuals to retrain in the genetic programming population in every generation, by default value is 1
5     verbosity=2, #it will state how much info TPOT will communicate while it is running
6     scoring='roc_auc', #use to evaluate the quality of given pipeline
7     random_state=42,
8     disable_update_check=True,
9     config_dict='TPOT light'
10 )
11 tpot.fit(X_train, y_train)
12
13 # AUC score for tpot model
14 tpot_auc_score = roc_auc_score(y_test, tpot.predict_proba(X_test)[:, 1])
15 print(f'\nAUC score: {tpot_auc_score:.4f}')
16
17 # Print best pipeline steps
18 print('\nBest pipeline steps:', end='\n')
19 for idx, (name, transform) in enumerate(tpot.fitted_pipeline_.steps, start=1):
20     # Print idx and transform
21     print(f'{idx}. {transform}')

```

```

Optimization Progress: 33%                                40/120 [00:06<00:21, 3.71pipeline/s]
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_split.py:700: UserWarning: The least populated class in y has only 1
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_split.py:700: UserWarning: The least populated class in y has only 1
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_split.py:700: UserWarning: The least populated class in y has only 1
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_split.py:700: UserWarning: The least populated class in y has only 1
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_split.py:700: UserWarning: The least populated class in y has only 1
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_split.py:700: UserWarning: The least populated class in y has only 1
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_split.py:700: UserWarning: The least populated class in y has only 1
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_split.py:700: UserWarning: The least populated class in y has only 1
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_split.py:700: UserWarning: The least populated class in y has only 1
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_split.py:700: UserWarning: The least populated class in y has only 1
warnings.warn(

```

```

-----
IndexError                                Traceback (most recent call last)
/usr/local/lib/python3.10/dist-packages/tpot/base.py in fit(self, features, target, sample_weight, groups)
    816         warnings.simplefilter("ignore")
--> 817         self._pop, _ = eaMuPlusLambda(
    818             population=self._pop,

```

⬆ 26 frames

IndexError: tuple index out of range

During handling of the above exception, another exception occurred:

```

ValueError                                Traceback (most recent call last)
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_ranking.py in binary_roc_auc_score(y_true, y_score, sample_weight, max_fpr)
1 tpot.fitted_pipeline_
--> 339         raise ValueError(

```

Model 1: Decision tree classifier

```

1 from sklearn.tree import DecisionTreeClassifier
2
3 clf = DecisionTreeClassifier(criterion='entropy', max_depth=8,
4                             min_samples_leaf=10,
5                             min_samples_split=6,
6                             random_state=42)
7 clf.fit(X_train,y_train)

```

```

▼
DecisionTreeClassifier
DecisionTreeClassifier(criterion='entropy', max_depth=8, min_samples_leaf=10,
min_samples_split=6, random_state=42)

```

```
1 y_predicted = clf.predict(X_test)
```

```
1 y_predicted
```

```
array([0, 0, 0, ..., 0, 0, 0])
```

```

1
2 print( classification_report(y_test, y_predicted) )

```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	1580
accuracy			1.00	1580
macro avg	1.00	1.00	1.00	1580
weighted avg	1.00	1.00	1.00	1580

```
1 accuracy_score(y_test,y_predicted)*100
```

```
100.0
```

```
1 #Creating Pipeline to see which model has more accuracy

1 from sklearn.preprocessing import StandardScaler
2 from sklearn.decomposition import PCA
3 from sklearn.pipeline import Pipeline
4 from sklearn.linear_model import LogisticRegression
5 from sklearn.tree import DecisionTreeClassifier
6 from sklearn.ensemble import RandomForestClassifier

1 pipeline_lr = Pipeline([('scaler1',StandardScaler()),
2                          ('pca1',PCA(n_components=2)),
3                          ('lr_classifier',LogisticRegression(random_state=0))])

1 pipeline_dt = Pipeline([('scaler2',StandardScaler()),
2                          ('pca2',PCA(n_components=2)),
3                          ('dt_classifier',DecisionTreeClassifier())])

1 pipeline_randomforest = Pipeline([('scaler3',StandardScaler()),
2                                    ('pca3',PCA(n_components=2)),
3                                    ('rf_classifier',RandomForestClassifier())])

1 pipeline = [pipeline_lr,pipeline_dt,pipeline_randomforest]

1 best_accuracy=0.0
2 best_classifier=0
3 best_pipeline=""

1 pipe_dict = {0:'Logistic Regression', 1:'Decision Tree', 2:'RandomForest'}
2
3 for pipe in pipeline:
4     pipe.fit(X_train,y_train)
5

1 for i,model in enumerate(pipeline):
2     print("{}Test Accuracy: {}".format(pipe_dict[i],model.score(X_test,y_test)))

Logistic RegressionTest Accuracy: 1.0
Decision TreeTest Accuracy: 0.9987341772151899
RandomForestTest Accuracy: 1.0
```

✓ 21s completed at 5:12 PM